

A **database instance** refers to the actual, real-time content of the data stored within a database at any specific point. It represents the current state of the database as opposed to the database schema, which is its structural definition.

### Key Characteristics of a Database Instance

1. **Real Data:** It holds the actual data in tables, such as customer details, product information, and orders.
2. **Dynamic:** Data changes over time as users add, update, or delete records.
3. **Snapshot:** It provides a snapshot of the database at a given time.
4. **Multiple Instances:** The same database schema can have different instances over time or across environments.

### Example: E-commerce Database

Let's consider a small-scale e-commerce system to illustrate the concept of a database instance.

#### Schema Definition:

- . **customers:** Stores customer information.
  - a. customer\_id, name, email
- . **products:** Stores product information.
  - a. product\_id, name, price, stock\_quantity
- . **orders:** Stores details of customer orders.
  - a. order\_id, customer\_id, order\_date, total\_amount

### Example Database Instance

Table: customers

A database instance holds the structural definitions of tables and relationships.

A database instance contains actual data that changes over time.

A database instance represents a snapshot of data at a specific point in time.

Multiple database instances can exist for the same database schema across different environments or timeframes.

A database instance defines the layout and constraints of tables without holding any data.

The database schema and instance are the same; one cannot change without affecting the other.

**1.6.3. Select the correct answer**

00:04 AA ☙ ☰ ☰

A company's HR system maintains a database with the following schema:

- Employee(EmpID, Name, DepartmentID, Salary)
- Department(DepartmentID, DeptName, Location)

At a particular point in time, the database contains the following data:  
Employee Table:

| EmpID | Name   | DepartmentID | Salary |
|-------|--------|--------------|--------|
| 101   | Rahul  | 1            | 50000  |
| 102   | Priya  | 2            | 60000  |
| 103   | Vikram | 1            | 55000  |

Department Table:

| DepartmentID | DeptName | Location  |
|--------------|----------|-----------|
| 1            | HR       | Mumbai    |
| 2            | IT       | Bangalore |

What is the best description of the data in the Employee and Department tables?

 Database Schema Normalized Schema Physical Schema Logical Schema Database Instance



### 1.7.1. Introduction to Database Design

00:05

Database design is the process of creating the structure and relationships for a database, ensuring efficient storage, management, and retrieval of data. A well-designed database ensures data accuracy, consistency, integrity, and optimal performance.

#### Goals of Database Design

- **Data Integrity:** Ensures accuracy and consistency by defining proper constraints (e.g., enforcing a valid email format).
- **Efficiency:** Optimises data storage and retrieval (e.g., using indexes for faster search results).
- **Flexibility:** Allows easy adaptation to future changes (e.g., adding new product categories in an e-commerce database).
- **Security:** Protects data with access control and encryption (e.g., restricting access to user passwords).
- **Usability:** Makes the database easy to use (e.g., clear table names and relationships for easier query writing).

#### Importance of Database Design

- **Data Organisation:** Makes it easier to manage and retrieve information (e.g., organising customer records by unique IDs).
- **Data Accuracy:** Reduces errors by enforcing rules (e.g., preventing entry of negative prices).
- **Performance Optimisation:** Ensures fast queries (e.g., using indexing on frequently queried columns).
- **Adaptability:** Easily accommodates future changes (e.g., adding new fields like customer preferences).
- **Data Security:** Ensures sensitive data is protected (e.g., encrypting personal customer information).
- **Maintenance Ease:** Simplifies database management and troubleshooting (e.g., easier

Use indexing on frequently searched columns, such as product categories, to enhance search speed.

Store customer passwords as plain text to simplify access for administrators.

Encrypt sensitive data such as payment details to enhance security.

Use clear and descriptive table names for better usability and maintenance.

Allow negative pricing for flexibility in special promotions.



## 1.7.2. Database Design: Key Terms

Imagine you're organising a digital library. How do you ensure every book, author, and borrower is tracked efficiently without confusion or duplication? That's where relational database concepts come in! In this section, you'll explore key terms like tables, rows, columns, keys, and more—each explained with real-world examples to help you design databases that are accurate, efficient, and easy to manage.

### Table

- **Definition:** Stores data in a structured format with rows and columns.
- **Usage:** Organises data into categories for easy access. Each row in a table represents a single record, while each column represents an attribute or field.
- **For example** let's take a "Books" table to catalog the library's collection:

| book_id | title                 | genre   |
|---------|-----------------------|---------|
| 1       | Harry Potter          | Fantasy |
| 2       | To Kill a Mockingbird | Classic |

The above table stores information about all the books in the library, including their IDs, titles, and genres.

### Row (Record)

- **Definition:** Represents a single entry in a table.
- **Usage:** Each row holds specific data about an individual entity.
- **For Example** let's take each row in the "Members" table represents a unique library member:

| member_id | name     | contact     |
|-----------|----------|-------------|
| 1         | John Doe | 123-4567890 |

Each row in the table represents a single book, and each column represents an attribute of the book.

- Each column in the table represents a single book, and each row represents an attribute of the book.
- The book\_id column is not necessarily, as the title and genre columns already provide enough information about each book.

- The title column represents the records in the table, while the book\_id and genre columns define the attributes.





### 1.7.3. Primary & Foreign Keys

#### Primary Key

- **Definition:** A unique identifier for each row in a table.
- **Usage:** Ensures no duplicate records and links tables.

#### Foreign Key

- **Definition:** A column in one table that links to the primary key of another table.
- **Usage:** Establishes relationships between tables.

Let's dive into the concepts of Primary Key and Foreign Key with a clear example using a library database.

#### 1. Books Table (Contains information about the books in the library)

Books Table

| book_id (Primary Key) | title                 | genre   |
|-----------------------|-----------------------|---------|
| 1                     | Harry Potter          | Fantasy |
| 2                     | To Kill a Mockingbird | Classic |

The book\_id column in the Books table is the primary key and uniquely identifies each book.

The member\_id column in the Loans table is a foreign key that links to the member\_id in the Members table, indicating which member borrowed a book.

The loan\_id column in the Loans table is a foreign key linking to the book\_id in the Books table, indicating which book was borrowed.

The book\_id column in the Books table is the primary key. It uniquely identifies each book in the library. No two books can have the same book\_id, ensuring each record in the table is unique.

#### 2. Members Table (Contains information about library members)



## 1.7.4. Key Terms - Contd

### Normalization

- **Definition:** Normalisation is the process of organising a database to avoid repeating the same information in multiple places.

- **Example:** Imagine you're running an online store. Instead of saving the customer's name, address, and phone number every time they make an order, you store their details in a separate "customers" table. Each order in the "orders" table then refers to the customer by their unique ID. This way, customer information isn't repeated with every order.

### Denormalization

- **Definition:** Denormalization means intentionally adding duplicate data in a database to make data retrieval faster. Normalisation can speed up data retrieval but may lead to data integrity challenges.
- **Example:** If you run a store and often need to calculate a customer's total spending, you might store the total amount in a "customers" table instead of calculating it every time from their orders. While this makes the process faster, it can cause issues if the total isn't updated correctly.

### Index

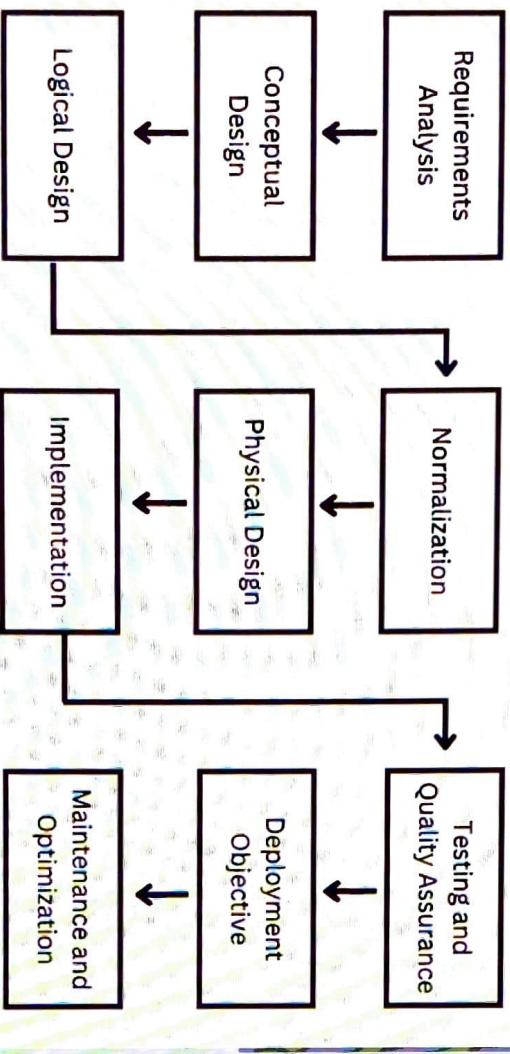
- **Definition:** An index is a tool that speeds up the search process in a database. Indexes help optimise query performance by providing a quick way to locate specific data within a table.
- **Example:** Think of a large library. If the books are arranged alphabetically by title, you can quickly find the book you're looking for. Similarly, in a database, creating an index on a column (like "product\_id" in a store) allows you to quickly search for specific products by their ID.

### Constraint

Normalization reduces data redundancy, while denormalization improves query performance by adding redundancy.

- Indexes are used to slow down database queries, especially when the data size is large.
- A constraint in a database is used to duplicate data for faster access.
- A schema is an actual database that stores user data, while the index is not needed for large databases.

Database design is a systematic process that plays a vital role in creating efficient, well-structured databases. A well-designed database ensures accurate and secure data storage, management, and retrieval. The design process involves several detailed steps, each contributing to a robust database system. Below is an overview of the major steps involved:



In the Analysis phase, interviewing only the head librarian would be sufficient since they manage the library operations.

During Conceptual Design, the ERD should show entities like 'Book', 'Member', and 'Loan' along with their relationships, such as "Member borrows Book."

Normalization is optional since the library data is straightforward and doesn't need to be broken down into smaller tables.

### 1. Analysis

**Objective:** Understand the specific needs and requirements of the database.

#### Activities:

- Conduct interviews with stakeholders (e.g., users, managers) to gather data needs.
- Document the types of data the database will store (e.g., customer details, orders).
- Identify data sources and formats (e.g., system logs, spreadsheets).

**Outcome:** A clear understanding of the database's purpose, data to be stored, and how it will



## 1.7.6. Database Design - Contd

00:05 AA ⚡ 🔍 🔍 ⌂

### 5. Physical Design

**Objective:** Translate the logical design into a format that is optimized for storage and performance on a specific DBMS.

#### Activities:

- Choose appropriate storage structures (e.g., indexes for faster searches, views for simplified queries).
- Define indexing strategies to speed up data retrieval.
- Optimize storage parameters for performance (e.g., data compression or partitioning).

**Outcome:** A physical schema that outlines how the data will be stored and accessed on the DBMS.

### 6. Implementation

**Objective:** Build the database system based on the logical and physical designs.

#### Activities:

- Write and execute SQL scripts to create tables, constraints, and other database objects.
- Populate the database with initial data.

**Outcome:** A fully functional database system ready for use.

### 7. Testing and Quality Assurance

**Objective:** Ensure the database functions correctly and meets the specified requirements.

#### Activities:

- Perform unit testing (testing individual components), integration testing (testing interactions between components), and user acceptance testing (ensuring the system meets user needs).
- Verify data accuracy, consistency, and reliability.

**Outcome:** A validated database that performs as expected and meets user requirements.

### 8. Deployment

Interviewing the bookstore's sales team to understand the types of data the system will need, such as book details, customer information, and order histories.

Writing SQL scripts to create tables for storing book, customer, and order data.

Creating an Entity-Relationship Diagram (ERD) to visualize the relationships between books, customers, and orders.

Allowing customers to access the database directly through a website for placing orders.

You are designing a military management system with the following tables. These tables ensure that soldiers, weapons, and assignments are tracked using primary and foreign keys.

### Weapons Table

| weapon_id (PK) | name  | type  |
|----------------|-------|-------|
| 1              | AK-47 | Rifle |
| 2              | M16   | Rifle |

### Soldiers Table

| soldier_id (PK) | name       | rank     |
|-----------------|------------|----------|
| 101             | John Doe   | Sergeant |
| 102             | Jane Smith | Private  |

### Assignments Table

| assignment_id (PK) | weapon_id (FK) | soldier_id (FK) | assignment_date | return_date |
|--------------------|----------------|-----------------|-----------------|-------------|
| 201                | 1              | 101             | 2024-12-01      | 2024-12-15  |

- The assignment\_id in the Assignments table is a foreign key that links to the weapon\_id in the Weapons table, ensuring each assignment is linked to a unique weapon.
- The weapon\_id in the Assignments table is a primary key, meaning it uniquely identifies each assignment and links to the Weapons table.

The soldier\_id in the Assignments table is a foreign key that links to the Soldiers table, showing which soldier is assigned a weapon, while weapon\_id is a foreign key linking to the Weapons table.

The Assignments table should have its own primary key for each soldier\_id and weapon\_id pair, eliminating the need for the assignment\_id as the primary key.

1.7.8. Select the correct answer

Which of the following steps in the database design process is focused on organising data to reduce redundancy and prevent data anomalies?

00:04 AA ⌂ ⌂ ⌂ ⌂

Analysis

Conceptual Design

Normalization

Testing and Quality Assurance

### 1.8.1. Types of Database Model



A database model defines the logical design and structure of a database, outlining how data is stored, accessed, and managed within a database management system (DBMS). It serves as a blueprint for organising and controlling data efficiently. In this tutorial, we will explore seven popular database models, each with distinct features and specific use cases.

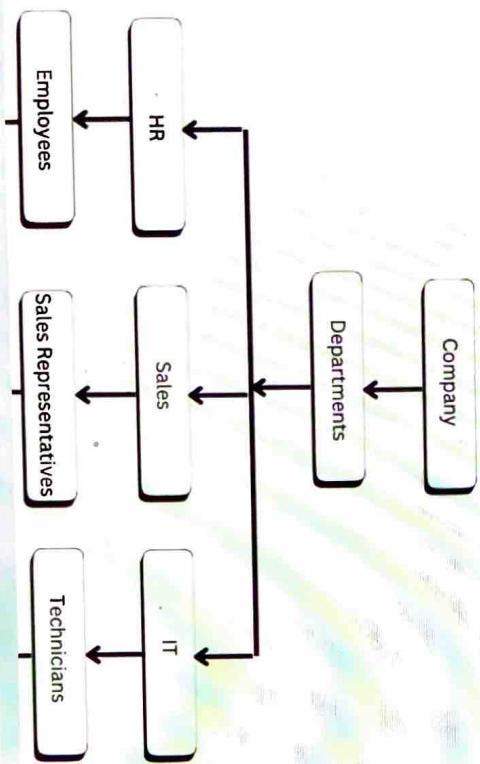
#### Hierarchical Model

- Data arranged in a tree-like structure with a single root.

- One-to-many relationships.
- Advantages: Fast data retrieval.

- Disadvantages: Limited flexibility, no many-to-many relationships.

In the below example, a company is the root, with departments like HR, Sales, and IT as branches. Each department has its own data, such as employees or sales representatives. This model works well for representing a structured organisation, like a company's hierarchy.



In the Hierarchical Model, data is arranged in a tree-like structure with a single root, and one-to-many relationships are supported.

The Graph Model is most suitable for complex data with intricate relationships, like social networks, where nodes represent entities and edges represent relationships.

The Graph Model cannot handle many-to-many relationships effectively.

The Relational Model simplifies data design by organizing it into tables and supporting SQL queries.



## 1.8.2. Types of Database Model- Contd

### Network Model

- Data in a graph-like structure with multiple parent nodes for a single child.
- Handles complex, many-to-many relationships.
- **Advantages:** Supports complex relationships and offers more flexibility.
- In this model, multiple entities (e.g., Departments, Employees, Projects) are connected to the company in a graph-like structure, with many-to-many relationships between entities.

### Object-oriented Model

- Data stored as objects, following object-oriented programming principles.
- **Advantages:** Supports complex structures and object-oriented features (e.g., MongoDB).

### NoSQL Model

- The NoSQL model is a type of database that stores unstructured or semi-structured data in formats like key-value pairs, documents, or graphs, offering flexibility, scalability, and high performance for large or complex datasets. (e.g., JSON, key-value pairs).
- Ideal for large data volumes and real-time analytics.
- **Advantages:** Scalable, high performance.

### Use Cases:

- The Hierarchical Model is suitable for a real-time scenario like an organisational hierarchy where a company has departments, and each department has employees.
- Relational Model is suitable for an online banking system where data like customers, accounts, and transactions are stored in tables with defined relationships.
- ER Model is suitable for designing the database schema for a hospital management system, representing entities like patients, doctors, and appointments.

The Network Model is designed to handle complex, many-to-many relationships and is ideal for applications like airline flight systems where multiple flights connect cities.

The Object-Oriented Model stores data as objects and is suitable for systems that require handling large volumes of unstructured data, such as social networks.

The NoSQL Model is optimized for large data volumes and real-time analytics, often storing data as JSON or key-value pairs for scalability and high performance.

The Network Model offers limited flexibility and is not well-suited for complex relationships like social networks.



## 1.8.3. Select the correct answer

00:04 A ⏪ ⏴ ⏵ ⏷

You are tasked with designing a database for a smart city management system that collects data from sensors across the city to monitor traffic, energy usage, and pollution levels in real time. The system also tracks citizen complaints and provides smart services like automated traffic light control and energy optimisation. Which of the following database models would be the most suitable for this use case?

- Hierarchical Model
- Entity-Relationship Model (ER)
- Relational Model
- Graph Model



## 1.9.1. Introduction to the ER Model

00:06

AA



A list of all registered customers in the system.

The Entity-Relationship (ER) model, introduced by Peter Chen in 1976, is fundamental for representing data and relationships in databases. It uses entities, attributes, and relationships to create a clear structure for database management.

In this section, we'll break down how the ER model works in the context of an e-commerce system, providing examples of real-world applications.

### 1. Entity

Entities are key objects or concepts within the system. They represent real-world elements that need to be stored in the database.

#### Example Entities in E-Commerce:

- **Customer:** Represents the person purchasing goods or services.
- **Product:** Represents an item available for purchase.
- **Order:** Represents a customer's order.

- **Payment:** Represents the transaction details for an order.

Each of these entities captures a critical part of the e-commerce system, from the customer making purchases to the payment method used for orders.

### 2. Entity Set

An entity set is a collection of entities of the same type. In simpler terms, it's a group of entities that belong to the same category.

#### Example Entity Sets:

- **Customer Set:** All registered customers in the system.
- **Product Set:** All available products in the store.
- **Order Set:** All orders placed by customers.
- **Payment Set:** All payment transactions made by customers.

Entity sets group similar data for easy management. For instance, the Customer Set includes all customers who have accounts on the platform.

- The attributes of the Product entity, such as product\_name and price.
- The relationship between a Customer and their Orders.
- The primary key used to uniquely identify an entity.

&lt; Prev

Reset

Submit

Next &gt;





## 1.9.2. Types of Attributes and Keys in the ER Model (Part 1)

00:07 AA ⏪ 🔍 ⏴ ⏵

Let's dive into attributes and keys in the ER model for an e-commerce system.

### 4. Types of Attributes

Attributes describe properties of entities. They come in various forms, each serving a different purpose:

- **Simple Attribute:** A basic, indivisible attribute.
- **Composite Attribute:** An attribute made up of smaller sub-attributes.
- 1. Example: address may include street, city, state, and zip\_code.
- **Derived Attribute:** An attribute whose value is derived from other attributes.
- 1. Example: total\_price is calculated based on the products in an order.
- **Multi-valued Attribute:** An attribute that can have multiple values.
- 1. Example: phone\_numbers can store multiple contact numbers for one customer.

### 5. Keys in the ER Model

Keys are used to uniquely identify entities. Here's a breakdown:

- **Primary Key:** An attribute or combination of attributes that uniquely identifies an entity within its set.

1. Example: In the Customer Set, customer\_id would be the primary key.

- **Candidate Key:** An alternative key that could also serve as the primary key. It uniquely identifies entities but may not be selected as the main identifier.

1. Example: The combination of customer\_email and customer\_phone might also serve as a candidate key for identifying a customer.

| customer_id | customer_name | customer_email    | customer_phone |
|-------------|---------------|-------------------|----------------|
| 1           | Alice         | alice@example.com | 1234567890     |

A simple attribute is an indivisible attribute and cannot be broken down further.

A derived attribute is an attribute whose value is explicitly stored in the database and cannot be calculated.

A multi-valued attribute allows multiple values for a single entity, such as multiple phone numbers for one customer.

A primary key is an attribute that uniquely identifies an entity and can never be composed of multiple attributes.

A foreign key refers to an attribute in one table that links to the primary key of another table, establishing a relationship between them.

< Prev Reset Submit Next >





### 1.9.3. Types of Keys and Relationships in the ER Model (Part 2)

00:07 AA ⏪ ☰

Now, let's look at the relationships in the ER model.

#### 7. Types of Relationships

Relationships in the ER model define how entities interact with each other. These relationships help map the connections and associations between different entities in the e-commerce system.

- **One-to-One (1:1):** Each entity in one set is related to exactly one entity in another. (one shipping address per customer).

- **One-to-Many (1:N):** One entity in the first set relates to multiple entities in the second set.

1. **Example:** One customer can place multiple orders (a customer can make many orders, but each order belongs to one customer).

1. **Example:** Multiple entities in one set relate to one entity in another.

- **Many-to-Many (N:N):** Entities in both sets are related to multiple entities in the other set.

1. **Example:** Many customers can add many products to their shopping cart (customers and products are related in a many-to-many fashion).

Consider a Library Management System that manages information about Books, Members, and Book Loans. The system includes the following tables:

**Books Table:** Contains information about books in the library.

|         |                                    |
|---------|------------------------------------|
| book_id | A unique identifier for each book. |
|---------|------------------------------------|

book\_id is the Primary Key in the Books Table, uniquely identifying each book.

isbn is a Candidate Key in the Books Table because it can uniquely identify a book, even though book\_id is used as the Primary Key.

member\_phone is a Composite Attribute because it can hold multiple contact numbers for a member.

loan\_id is the Primary Key in the Book Loans Table, uniquely identifying each book loan.

book\_title is a Derived Attribute because it is calculated based on the book's isbn.

loan\_date is a Simple Attribute because it stores a single piece of data: the date when the book was borrowed.



#### 1.9.4. Components of ER Diagram: Entities

An Entity Relationship (ER) Diagram in DBMS is a blueprint that represents the database structure, which can later be implemented as tables. It helps systematically study data requirements before creating the actual database.

The ER Model models the logical view of the system and uses these symbols:

- Rectangles:** Rectangles represent Entities in the ER Model.
- Ellipses:** Ellipses represent Attributes in the ER Model.
- Diamonds:** Diamonds represent Relationships among Entities.
- Lines:** Lines represent attributes to entities and entity sets with other relationship types.
- Double Ellipses:** Double Ellipses represent Multi-Valued Attributes.
- Double Rectangles:** Double Rectangle represents a Weak Entity.

| Shape          | Symbol  | Represents  |
|----------------|---|---|
| Rectangle      |  | Entities in ER Model Diagram  |
| Ellipse        |  | Attributes to an Entity   |
| Diamond        |  | Relationship among Entities   |
| Line           |  | Connects Attributes to Entities<br>Connects Entity Sets with other Relation Types |
| Double Ellipse |  | Multi-Valued Attributes   |

1 : A, 2 : B, 3 : C, 4 : D

1 : D, 2 : C, 3 : B, 4 : A

1 : B, 2 : A, 3 : D, 4 : C

[< Prev](#) [Reset](#) [Submit](#) [Next >](#)





## Attributes

Attributes are the properties or characteristics that define an entity type. These attributes describe various aspects of the entities.

LYMPS OF LARINUS.

- called the key attribute.

- **Composite attribute:** An attribute composed of many other attributes is called a composite attribute.

- a. **Example:** The "shipping address" attribute is composed of sub-attributes such as "country", "state", "city" and "street".**Representation:** An oval comprising smaller ovals.

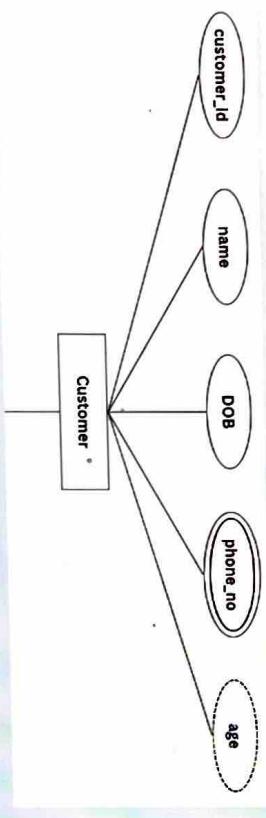
- **Multivalued attribute:** An attribute consisting of more than one value for a given entity.

a. **Example:** The "phone numbers" attribute can hold more than one phone number  
*for a single customer entity*

**Representation:** A double oval.

- **Derived attribute:** An attribute that can be derived from other attributes of the entity type is known as a derived attribute.

- Attribute Representation:** A dashed oval.



- A Strong Entity in the ER Model does not rely on other entities and has a unique identifier.

- Weak Entities can have key attributes independently

- Multivalued attributes in the ER Model are represented by rectangles.

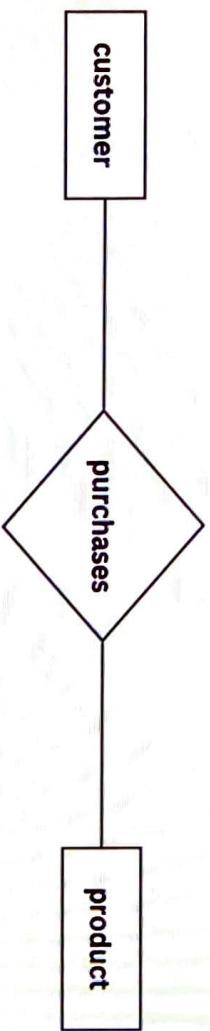


### 1.9.6. Components of ER Diagram: Relationship Type and Relationship Set

00:05

 1 : B, 2 : A, 3 : C

A Relationship Type represents the association between entity types.  
For Example:  
**Purchase Relationship Type:** This relationship type exists between the "customer" entity type and the "product" entity type. It signifies that a customer can purchase products from the platform.



#### Entity-Relationship Set:

An entity-relationship set is a collection of relationships of the same type.

#### Purchase Relationship Set:

This set includes instances of the "purchase" relationship type. Within this set, you might have instances where "customers" entity instances (e.g., C1, C2) are associated with "products" entity instances (e.g., P1, P2) to represent actual purchases.

**Degree of a Relationship Set:** In the realm of database design, the degree of a relationship set refers to the count of distinct entity sets that take part in that specific relationship set.

**1. Unary Relationship:** A unary relationship is characterised by the involvement of a single entity set within the relationship.

- Example:** Customer Follows Customer: customers can follow other customers to get updates on their shopping activities or product reviews. Each customer is related to another customer through the "follows" relationship. It's a unary relationship because only the "customer" entity set is involved.



## 1.9.7. Cardinality & Participation Constraint

00:04 AA ☺ ☐ ☒ ☓ ☔

### Cardinality

It refers to the number of times an entity in an entity set participates in a relationship set. Cardinality plays a crucial role in defining how entities are related to each other within the database.

**1. One-to-One (1:1):** In a one-to-one relationship, each entity in each entity set can participate only once in the relationship.

- **Example:** One customer can have only one active shipping address, and one active shipping address is associated with only one customer.
- **Representation:** One-to-One Cardinality (1:1)



**2. One-to-Many (1:N):** In a one-to-many relationship, entities in one entity set can participate once, but entities in the other entity set can participate more than once in the relationship.

- **Example:** A customer can place multiple orders, but each order is associated with only one customer.
- **Representation:** One-to-Many Cardinality (1:N)



In a one-to-many relationship, entities in one entity set can participate more than once, while entities in the other entity set can participate only once.

Total participation is represented by a double line in the ER diagram, indicating that each entity in the entity set must participate in the relationship.

In a many-to-many relationship, entities in all entity sets can participate only once in the relationship.

Partial participation implies that each entity in the entity set must participate in the relationship.



## 1.9.8. Generalization, Specialization and Aggregation

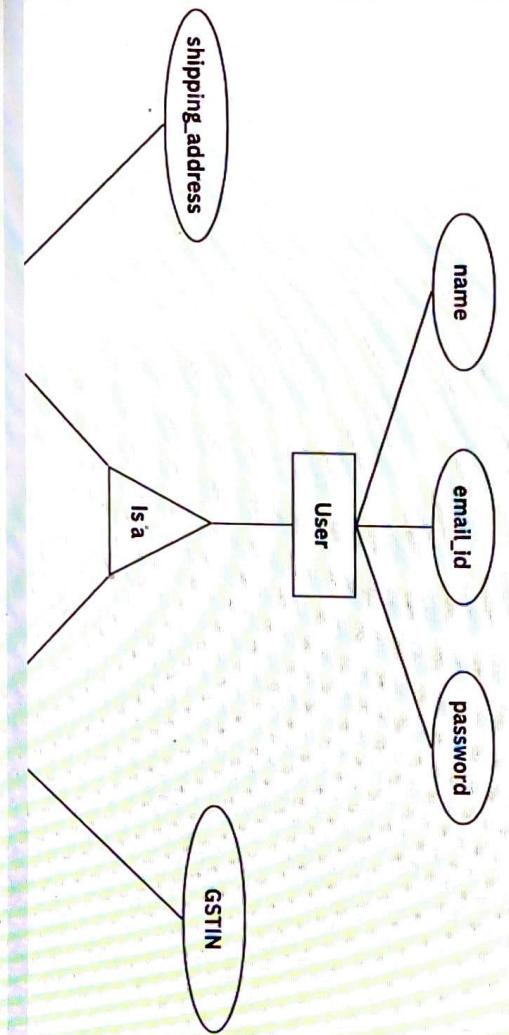
00:06 A ⏪ ⏴ ⏵ ⏷

In the Entity-Relationship (ER) model, Generalization, Specialization, and Aggregation are important concepts used to represent complex relationships and hierarchy in a database system. These concepts provide a way to organize and model data more effectively.

### Generalization

Generalization is a top-down approach used to represent a higher-level entity by combining lower-level entities that share common characteristics. It is a process of abstracting common properties from multiple entity types to create a more general entity type. The resulting entity is known as a superclass or parent entity, while the original entities are called subclasses or child entities.

**Example:** You can have a general "User" entity that encompasses "customer" and "seller" as subclasses. The "user" entity may have common attributes like username and password, while the "customer" and "seller" entities have additional attributes specific to their roles.



Generalization involves abstracting common properties from multiple entity types (subclasses) to create a superclass (parent entity).

Specialization divides a higher-level entity into lower-level entities based on unique attributes.

- Aggregation allows you to treat relationships between entities as separate entities.
- In a Generalization hierarchy, the parent entity is called a subclass.
- Aggregated entities are typically represented as rectangles in an ER diagram.



### 1.10.1. The Relational Model: A Game-Changer in Data Management

00:06 AA ⏪ ⏴ ⏵ ⏷

The Relational Model, introduced by Dr. Edgar F. Codd in the early 1970s at IBM, revolutionized how we store and retrieve data. His groundbreaking paper, "A Relational Model of Data for Large Shared Data Banks," laid the foundation for modern databases, making data management more logical and efficient.

### Key Concepts of the Relational Model

#### Employee Table (Relation)

| Primary Key → | EmployeeID | Name  | Salary |
|---------------|------------|-------|--------|
| 1120          | Vaibhav    | 51000 |        |
| 1231          | Neha       | 53000 |        |
| 1235          | Harsh      | 45000 |        |
| 1250          | Shreya     | 65000 |        |

Columns (Attributes)

Tuples (Rows)

- It is a unique identifier for each record in a table.
- It is used to enforce data integrity within a single table.
- It defines the attributes or columns of a table.
- It establishes a link between two tables by referencing a primary key in another table.

1. **Tables (Relations):** Data is organized into tables, with rows and columns resembling a spreadsheet. Each table represents an entity, like "Customers" or "Products."
2. **Rows (Tuples):** Each row represents a single record, such as a customer or product.
3. **Columns (Attributes):** Columns define the type of data stored, such as Name or Price.
4. **Relationships:** Tables are linked using Primary Keys (unique identifiers) and Foreign Keys (non-primary fields) allowing data to be related across tables.

## Relational Model Components

### 1. Tables (Relations)

- Tables are the backbone of the Relational Model, organizing data into structured entities. Each table represents a specific type of data, consisting of rows (records) and columns (attributes).

- Example Table (Products):

| product_id | product_name | price |
|------------|--------------|-------|
| 1          | Laptop       | 800   |
| 2          | Smartphone   | 500   |
| 3          | Headphones   | 100   |

### 2. Rows (Tuples)

- Rows represent individual records in a table. Each row contains data values for the table's attributes.

- Example Table (E-commerce Database):

| product_id | product_name | price |
|------------|--------------|-------|
| 1          | Laptop       | 800   |

← row

The Customer\_ID in the Orders table is a Primary Key.

The Customer\_ID in the Orders table is a Foreign Key referencing the Customer\_ID in the Customers table.

There is a Many-to-Many relationship between the Customers and Orders tables.

The Orders table demonstrates a One-to-Many relationship with the Customers table.

Each customer can place multiple orders.



### 1.10.3. Advantages and Features of the Relational Model

00:07 AA ☺ ☐ ☐

In this section, we explore the advantages and features of the relational model which is widely used in e-commerce to manage data efficiently, ensuring reliability, scalability, and ease of use.

#### Advantages

##### 1. Data Integrity:

- Primary Keys ensure each record is unique, and Foreign Keys maintain accurate relationships between tables, preventing invalid data entries.

- Example: product\_id in the "products" table guarantees every product is unique.

##### 2. Flexibility:

- The use of SQL enables powerful queries to retrieve and analyze data from multiple tables, making it easy to filter or join data.

- Example: Fetch customer orders by linking the "customers" and "orders" tables.

##### 3. Scalability:

- Relational databases grow with your business. Adding new tables (e.g., reviews or shipping details) does not affect existing data structures.

- Example: Expand to include "reviews" or "shipping" tables without disrupting existing data.

##### 4. Data Consistency:

- Through Normalization, relational databases reduce duplication, ensuring consistent data across tables. For instance, customer details are stored in one table instead of being repeated in every order.

- Example: Store customer details in the "customers" table, not in every order.

#### Features

- Tables: Data is organized into structured-tables representing entities like "customers" or "products."

- Relationships: Tables are linked using Primary Keys and Foreign Keys, enabling connections across data

 Primary Keys ensure data redundancy across tables for easy retrieval. Foreign Keys are used to link related data across multiple tables. SQL enables efficient insertion, retrieval, and updating of data in relational databases. Normalization reduces data redundancy and improves consistency. Relational databases cannot scale easily as the data grows.

## 1.10.4. Transforming an Entity-Relationship (ER) Model into a Relational M...

00:06



Converting an ER model to a Relational Model bridges the conceptual database design with its practical implementation in a relational database management system (RDBMS). Here's a concise, step-by-step guide:

### 1. Entities to Tables

Each entity in the ER model becomes a table in the relational model, with attributes as columns.

#### Example:

- The "Product" entity (product\_id, name, price) becomes the "Products" table.
- The "Customer" entity (customer\_id, first\_name, last\_name) becomes the "Customers" table.

### 2. Relationships to Foreign Keys

Relationships between entities are implemented using Foreign Keys:

#### Example:

- For a "Purchases" relationship between "Customers" and "Products," add customer\_id in the "Orders" table as a Foreign Key referencing customer\_id in "Customers".

### 3. Cardinality to Constraints

Translate cardinality (1:1, 1:N, M:N) into relational constraints:

- One-to-One (1:1):** Add a Foreign Key in either table.
- Example:** "Customer" and "Address" -> Add customer\_id to the "Addresses" table.

- One-to-Many (1:N):** Add a Foreign Key to the "N" side table.

- Example:** "Customer" and "Order" -> Add customer\_id in the "Orders" table.

- Many-to-Many (M:N):** Create a junction table with Foreign Keys referencing both entities.

- Example:** "Products" and "Categories" -> Create "Product\_Categories" table with

The customer\_id in the "Orders" table is a Foreign Key referencing the "Customers" table.

The product\_id in the "Products" table is both a Primary Key and a Foreign Key.

The "Product\_Categories" table must include at least two Foreign Keys referencing "Products" and "Categories."

The "Orders" table can exist without a relationship with the "Customers" table.

A many-to-many relationship requires a junction table to link the related entities.